

Roger Richard Demello

Nagpur, India

+91 94217 78898 | rogerdemello289@gmail.com | linkedin.com/in/rogerdemello | github.com/rogerdemello | rogerdemello.tech

Summary

Engineer building AI systems that hold up outside a demo - agents wired into real platforms, a recommender with a full MLOps loop, and evaluation harnesses that publish their own negative results. Electronics & Communication Engineering graduate with a Machine Learning minor. Shipped backend services behind **3 production applications** serving **200+ users**, cut response latency **35%** at **99.5% uptime**, and compressed release cycles **50%** through automated validation.

Technical Skills

Languages: Python, TypeScript / JavaScript, C++, C, SQL, Bash

AI & Machine Learning: LLM agents and multi-agent orchestration (LangGraph, MCP), RAG with source citations and vector search, PyTorch, scikit-learn, LightGBM, two-tower retrieval and ranking, anomaly detection, evaluation harnesses and published benchmarks

Backend & Data: FastAPI, Flask, Node.js / Express, Next.js / React, PostgreSQL / SQLite / Supabase, Redis, Kafka, REST APIs and webhooks, authentication and RBAC, audit logging, pandas / NumPy / SciPy

MLOps & Infrastructure: MLflow model registry, retraining and precompute jobs, Prometheus / Grafana, Docker, Kubernetes-ready deployments, Git, GitHub Actions CI, pytest, AWS (Certified Cloud Practitioner)

Blockchain & Web3: Solidity and Sui Move contracts, on-chain consent and audit trails, decentralized storage and encryption, wallet integration

Experience

AI LifeBOT

Jan 2026 - Jun 2026

AI Engineer Intern

Remote

- Built and shipped backend services behind **3 production applications** serving **200+ users**, covering REST endpoints, data ingestion, and storage integration.
- Debugged and optimised pipeline hot paths to cut response latency **35%** at **99.5% uptime**, and added telemetry plus automated test pipelines to catch regressions before release.
- Broke features into small deliverable tasks, wrote the accompanying documentation, and automated validation to compress release cycles **50%** while working across product and engineering.

Center for Manufacturing (CFM), RCOEM

May 2025 - Jul 2025

Machine Learning Research Intern

Nagpur, India

- Built data cleaning and preprocessing pipelines over **1,000+ records**, cutting pipeline development time **30%**, and presented both the results and their limitations to stakeholders.

Projects

ML Guardian | *Python, FastAPI, MCP, DataHub, Gemini* | [GitHub](#)

- Autonomous ML reliability agent on DataHub's context platform that catches silent production failures - stale upstreams, creeping null rates, renamed columns - before a KPI moves.
- Runs a scan → score → incident → write-back → remediate loop that names the downstream models and dashboards at risk and writes findings back as DataHub tags; **offline-first** fixtures run with no Docker or API keys.

music-recsys | *Python, PyTorch, LightGBM, MLflow, Prometheus* | [GitHub](#)

- Two-stage recommender - two-tower embeddings feeding ANN retrieval into a **LightGBM** ranker - with the full MLOps loop around it: event bus, feature updater, online store, model registry, retrain and precompute jobs.
- Put cache, event bus, store, ANN index and model registry behind **Protocol interfaces**, so CPU-only local implementations swap to Redis, Kafka and Kubernetes by changing one config value.

Autonomous Executive Email Copilot | *Python, FastAPI, Pydantic, SciPy, React, pytest* | [GitHub](#)

- Deterministic simulation environment scoring email-triage agents against **bounded, numerically stable** metrics across baseline, stress-test perturbation, LLM-driven and multi-agent hybrid policy modes.
- Shipped the operational stack around it - FastAPI backend, React dashboard, approval workflows, episode replay, telemetry - and published the benchmark results including the negative findings.

Engram | *Next.js, TypeScript, Sui Move, Walrus, Seal* | [GitHub](#) | [Live](#)

- User-owned, verifiable memory layer for AI agents on the **Sui** blockchain - on-chain consent grants and revokes, Seal-encrypted Walrus storage, and an audit trail of every memory access.
- Agents cite the exact memories they use, so users can audit and correct AI responses; demonstrated by a multi-app scenario granting and revoking consent on-chain in real time.

Education

Shri Ramdeobaba College of Engineering and Management (RCOEM)

2022 - 2026

B.Tech, Electronics & Communication Engineering

GPA: 8.90/10

Minor: Artificial Intelligence & Machine Learning

GPA: 9.67/10

Certifications & Awards

AWS Certified Cloud Practitioner (Oct 2025)

2nd Place - ByteSize Sage AI National Hackathon

Finalist - Paytm × Sarvam × Logitech AI National Hackathon